

Venti

Presented By:
陈子暘

Outline

Abstract
Background

Venti

Data
Organiza-
tion[Organization]

Application
Example
Vac

Venti: a new approach to archival storage¹

Presented By: 陈子暘

Fudan University

13307130148@fudan.edu.cn

January 2, 2015

¹powered by Xe₃La₂TeX

Outline

Venti

Presented By:
陈子暘

Outline

Abstract
Background

Venti

Data
Organization
[Organization]

Application
Example
Vac

- 1 Outline
 - Abstract
 - Background

- 2 Venti

- 3 Data Organization[Organization]

- 4 Application Example

Abstract

Venti

Presented By:

陈子暘

Outline

Abstract

Background

Venti

Data

Organization
[Organization]

Application

Example

Vac

- Venti: A network storage system intended for archival data
- A building block for a variety of storage applications
 - ① logical backup
 - ② physical backup
 - ③ snapshot file systems
- A block is identified by a unique hash of it's contents
- Enforce a write-once policy
- Duplicate copies of a block can be coalesced

Archival Storage

Venti

Presented By:

陈子暘

Outline

Abstract

Background

Venti

Data

Organization
[Organization]

Application

Example

Vac

- Purpose

- Store data for long periods of time (forever)
- Data may not be needed frequently, but when it is needed it is often crucial

Prevalent Form

Venti

Presented By:

陈子暘

Outline

Abstract

Background

Venti

Data

Organiza-
tion[Organization]

Application

Example

Vac

- Tape backup
 - Backup data to magnetic tape
 - (tar, ufsdump...)
 - Full backup vs Incremental backup
 - To provide backup as a central service for a number of client machines

Prevalent Form

Venti

Presented By:
陈子暘

Outline

Abstract

Background

Venti

Data

Organiza-
tion[Organization]

Application
Example

Vac

- Snapshot

- A snapshot is a consistent read-only view of the file system at some point in the past.
- Each snapshot is a complete file system tree, much like a full backup.
- A snapshot only requires additional storage for the blocks that have changed, like a incremental backup.
- Always available and easy to access
- Plan 9, WAFL, AFS...

Outline

Venti

Presented By:
陈子暘

Outline

Abstract
Background

Venti

Data
Organization
[Organization]

Application
Example

Vac

1 Outline

2 Venti

3 Data Organization[Organization]

4 Application Example

Venti Archival Storage

Venti

Presented By:

陈子暘

Outline

Abstract

Background

Venti

Data

Organization[Organization]

Application

Example

Vac

- Goal: To provide a write-once archival repository than can be shared by multiple client machines and applications.
- Block level network storage system
 - Actually a backend storage for client apps
- Blocks addressed by hash of their contents
 - Use SHA-1 algorithm
 - Use hash value as its unique 'fingerprint'
- Write-Once policy
 - Block once written, never modified
 - Modified blocks will have new address

Why SHA-1?

Venti

Presented By:
陈子阳

Outline

Abstract
Background

Venti

Data
Organization
[Organization]

Application
Example
Vac

- SHA-1 hash function is developed by NIST
- Output 160 bit hash values(20 bytes)
- Probability that there will be one or more collisions:

$$p \leq \frac{n(n-1)}{2} \times \frac{1}{2^b}$$

- Consider a large storage system contains 10^{18} byte of data stored as 8 Kbyte blocks($\sim 10^{14}$ blocks), the probability is less than 10^{-20} .
- Variants of SHA-1 can produce 256, 384 and 512 bit results for future use.

Venti Archival Storage

Venti

Presented By:

陈子暘

Outline

Abstract

Background

Venti

Data

Organization
[Organization]

Application
Example

Vac

- Multiple clients can Share a Venti server
 - Hash function gives an universal namespace
 - Duplication increases the utility rate of space
- Inherent integrity checking for data
- Caching is simplified
- Uses magnetic disk as storage technology
 - Access time comparable to non-archival data

Outline

Venti

Presented By:
陈子暘

Outline

Abstract
Background

Venti

Data
Organization
[Organization]

Application
Example
Vac

1 Outline

2 Venti

3 Data Organization[Organization]

4 Application Example

Data Organization

Venti

Presented By:

陈子暘

Outline

Abstract

Background

Venti

Data
Organiza-
tion[Organiza-

Application
Example

Vac

- Data is divided into blocks and written to the server
- Pack the fingerprints into additional blocks, called pointer blocks, that are also written to the server
- Until a single fingerprint is obtained
- Applications can use such a structure to store a single file or to mimic the behavior of a physical device such as a tape or a disk drive

Data Organization

Venti

Presented By:
陈子暘

Outline

Abstract
Background

Venti

Data
Organiza-
tion[Organiza-

Application
Example
Vac




Figure 1. A tree structure for storing a linear sequence of blocks

Data Organization

Venti

Presented By:
陈子暘

Outline

Abstract
Background

Venti

Data
Organiza-
tion[Organization]

Application
Example
Vac

- Venti does not allow such a tree to be modified
- But new versions of the tree can be generated efficiently by storing the new or modified data blocks and reusing the unchanged sections
- By mixing data and fingerprints in a block, more complex data structures can be constructed.
- For example, a structure for storing a file system may include three types of blocks:
 - Directory
 - Pointer
 - Data.

Data Organization

Venti

Presented By:
陈子旸

Outline

Abstract
Background

Venti

Data
Organiza-
tion[Organizatio

Application
Example
Vac




Figure 2. Build a new version of the tree.

Outline

Venti

Presented By:
陈子暘

Outline

Abstract
Background

Venti

Data
Organization
[Organization]

Application Example

Vac

1 Outline

2 Venti

3 Data Organization[Organization]

4 Application Example

- Vac

Vac

Venti

Presented By:
陈子暘

Outline

Abstract
Background

Venti

Data
Organization
[Organization]

Application
Example

Vac

- Vac is an application similar to tar and zip
 - With vac, Selected files will be stored as a tree of blocks on Venti server.
 - The output is always 45 bytes long, included a 20 byte root fingerprint.
 - 'unvac' enables user to restore files from a vac archive.
- Vac writes each file as a separated collection of Venti blocks, which can coalesce duplicate copies of a file
- Incremental backups options can improve performance